

第8课：彻底实战详解使用IDE开发Spark程序--集群模式运行

原创

2016年01月30日 08:41:27

标签：spark

1386

第8课：彻底实战详解使用IDE开发Spark程序--集群模式运行

拷贝WordCount.scala生成WordCountCluster.scala。

1. 将object WordCount改为object WordCountCluster

2. 将conf.setMaster("local")行注释掉。在提交时再配置。

3. 将文件源修改为val lines = sc.textFile("hdfs://192.168.1.121:9000/user/spark/README.md")

4.启动hadoop和spark，并启动spark的history-server.sh

5.导出jar包。在WordCount Project上点击右键->export，选择java下的JAR file，点击Next。选择好导出路径后点击finish。

6.将导出的jar包拷贝到虚拟机中/home/richard/spark-1.6.0/class目录下。

7. 使用spark-submit提交运行jar包。

sprak-submit --class com.dt.spark.WordCountCluster --master spark://slq1:7077 /home/richard/spark-1.6.0/class/WordCount.jar

集群模式代码：

package com.dt.spark

import org.apache.spark.SparkConf
import org.apache.spark.SparkContext

```
/**
 * 使用scala开发集群运行的spark WordCount程序
 * DT大数据梦工厂
 * 新浪微博：http://weibo.com/ilovepains/
 */
object WordCountCluster {
  def main(args: Array[String]){
    /**
     * 第一步：创建spark的配置对象SparkConf，设置Spark程序运行时的配置信息
     * 例如通过setMaster来设置程序要链接的Spark集群的Master的URL，如果设置
     * 为local，则代表Spark程序在本地运行，特别适合于机器配置条件非常差（例如
     * 只有1内存）的初学者
     */
    val conf = new SparkConf() //创建SparkConf对象。因为是全局唯一的，所以使用new，不用工厂方法模式。
    conf.setName("Wow, My First Spark App!") //设置应用程序的名称，在程序运行的监控界面可以看到名称
    // conf.setMaster("local") //此时程序在Spark集群。

    /**
     * 第二步：创建SparkContext对象，
     * SparkContext是Spark程序所有功能的唯一入口，无论是采用scala/java/Python/R等都必须有一个SparkContext，
     而且默认都只有一个。
     * SparkContext核心作用：初始化应用程序运行时所需要的核心组件，包括DAGScheduler,TaskScheduler,Scheduler
    Backend,
     * 同时还会负责Spark程序往Master注册程序等。SparkContext是整个Spark应用程序中最为重要的一个对象，
     *
     */
    val sc = new SparkContext(conf) //通过创建SparkContext对象，通过传入SparkConf实例来定制Spark地的具体参数和配置信息。
    /**
     * 第三步：根据具体的数据来源（HBase/Local FS/DB/S3等）通过SparkContext创建RDD，
     * RDD创建有三种基本方式：1.根据外部数据来源（如HDFS），2.根据Scala集合，3.由其他RDD操作产生
     * 数据会被RDD划分成为一系列的Partitions，分配到每个Partition的数据属于一个Task的处理范畴，
     */
    val lines = sc.textFile("hdfs://192.168.1.121:9000/user/spark/README.md") //读取HDFS文件，并切分成不同的partitions
    //也可以写成：lines:RDD[String] = sc.textFile 类型推断
    /**
     * 第4步：对初始RDD进行Transformation级别的处理。例如map/filter等高阶函数等的编程
     * 来进行具体的数据计算。第4.1步：将每一行的字符串拆分成单词的单词。
     */
    val words = lines.flatMap { line => line.split(" ") } //对每一行的字符串进行单词拆分，map每次循环一行，将每一行的小集合通过flat合并成一个大集合
    /**
     * 第4.2步，在单词拆分的基础上对每个单词实例 进行计数为1，也就是word => (word,1)
     */
    val pairs = words.map { word => (word,1) }
    /**
     * 第4.3步，在每个单词实例计数为1的基础上，统计每个单词在文件中出现的总次数。
     */
    val wordCounts = pairs.reduceByKey(_+_ ) //对相同的Key,进行Value的累计（包括Local和Reduce级别同时Reduce）
    wordCounts.collect.foreach(wordNumberPair => println(wordNumberPair._1 + " : " + wordNumberPair._2))
    sc.stop() //把上下文去掉，释放资源

  }
}
```

运行时的log：

```
[richard@slq1 bin]$ ./spark-submit --class com.dt.spark.WordCountCluster --master spark://slq1:7077 /home/richard/spark-1.6.0/class/WordCount.jar
16/01/30 08:16:06 INFO spark.SparkContext: Running Spark version 1.6.0
16/01/30 08:16:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
16/01/30 08:16:12 INFO spark.SecurityManager: Changing view acls to: richard
16/01/30 08:16:12 INFO spark.SecurityManager: Changing modify acls to: richard
16/01/30 08:16:12 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(richard); users with modify permissions: Set(richard)
16/01/30 08:16:19 INFO util.Utils: Successfully started service 'sparkDriver' on port 34985.
16/01/30 08:16:23 INFO slf4j.Slf4jLogger: Slf4jLogger started
16/01/30 08:16:24 INFO Remoting: Starting remoting
16/01/30 08:16:26 INFO Remoting: Remoting started; listening on addresses : [akka.tcp://sparkDriverActorSystem@192.168.1.121:33547]
16/01/30 08:16:26 INFO util.Utils: Successfully started service 'sparkDriverActorSystem' on port 33547.
16/01/30 08:16:26 INFO spark.SparkEnv: Registering MapOutputTracker
16/01/30 08:16:27 INFO spark.SparkEnv: Registering BlockManagerMaster
16/01/30 08:16:27 INFO storage.DiskBlockManager: Created local directory at /tmp/blockmgr-a020b1d8-f908-4473-852a-f6a55b545e02
16/01/30 08:16:27 INFO storage.MemoryStore: MemoryStore started with capacity 517.4 MB
16/01/30 08:16:28 INFO spark.SparkEnv: Registering OutputCommitCoordinator
16/01/30 08:16:31 INFO server.Server: jetty-8.y.z-SNAPSHOT
16/01/30 08:16:32 INFO server.AbstractConnector: Started SelectChannelConnector@0.0.0.0:4040
16/01/30 08:16:32 INFO util.Utils: Successfully started service 'SparkUI' on port 4040.
```



+ 关注

原创

粉丝

喜欢

码云

116

0

0

未开通

他的最新文章

[更多文章](#)

第72课：Spark UDF和UDAF解密学习笔记

第71课：Spark SQL窗口函数解密与实战学习笔记

第70课:SparkSQL内置函数解密与实战学习笔记

c++

游戏编程

小程序

加盟

多用户商城系统

移民澳洲的条件

网上商城系统

蓝鸟

新西兰移民条件

苹果6s降价

马来西亚房价

开源商城系统

上海青年公寓

二手钢琴

在线课程

免费直播 神经网络的原理及结构设计
讲师：何宇健Apache Weex：移动研发的进阶之路
讲师：董岩

热门文章

第70课:SparkSQL内置函数解密与实战学习笔记
4404第69课：SparkSQL通过Hive数据源实战学习笔记
3588第72课：Spark UDF和UDAF解密学习笔记
3472第64课：SparkSQL下Parquet的数据切分和压缩内幕详解学习笔记
3411第68课：SparkSQL JDBC实战详解学习笔记
3183

16/01/30 08:16:32 INFO ui.SparkUI: Started SparkUI at http://192.168.1.121:4040
16/01/30 08:16:32 INFO spark.HttpFileServer: HTTP File server directory is /tmp/spark-42bd08b-938a-45fe-8dae-81db69953855/httpd-ccd8a47c-1b7f-4ac6-b7e9-510b4f88c4b2
16/01/30 08:16:32 INFO spark.HttpServer: Starting HTTP Server
16/01/30 08:16:32 INFO server.Server: jetty-8.y.z-SNAPSHOT
16/01/30 08:16:32 INFO server.AbstractConnector: Started SocketConnector@0.0.0.0:50966
16/01/30 08:16:32 INFO util.Utils: Successfully started service 'HTTP file server' on port 50966.
16/01/30 08:16:33 INFO spark.SparkContext: Added JAR file:/home/richard/spark-1.6.0/class/WordCount.jar at http://192.168.1.121:50966/jars/WordCount.jar with timestamp 1454112993094
16/01/30 08:16:34 INFO client.AppClient\$ClientEndpoint: Connecting to master spark://slq1:7077...
16/01/30 08:16:37 INFO cluster.SparkDeploySchedulerBackend: Connected to Spark cluster with app ID app-20160130081636-0001
16/01/30 08:16:37 INFO client.AppClient\$ClientEndpoint: Executor added: app-20160130081636-0001/0 on worker-20160130074034-192.168.1.123-57185 (192.168.1.123:57185) with 1 cores
16/01/30 08:16:37 INFO cluster.SparkDeploySchedulerBackend: Granted executor ID app-20160130081636-0001/0 on hostPort 192.168.1.123:57185 with 1 cores, 1024.0 MB RAM
16/01/30 08:16:37 INFO client.AppClient\$ClientEndpoint: Executor added: app-20160130081636-0001/1 on worker-20160130074035-192.168.1.122-37406 (192.168.1.122:37406) with 1 cores
16/01/30 08:16:37 INFO cluster.SparkDeploySchedulerBackend: Granted executor ID app-20160130081636-0001/1 on hostPort 192.168.1.122:37406 with 1 cores, 1024.0 MB RAM
16/01/30 08:16:37 INFO client.AppClient\$ClientEndpoint: Executor added: app-20160130081636-0001/2 on worker-20160130074053-192.168.1.121-45928 (192.168.1.121:45928) with 1 cores
16/01/30 08:16:37 INFO cluster.SparkDeploySchedulerBackend: Granted executor ID app-20160130081636-0001/2 on hostPort 192.168.1.121:45928 with 1 cores, 1024.0 MB RAM
16/01/30 08:16:37 INFO util.Utils: Successfully started service
'org.apache.spark.network.netty.NettyBlockTransferService' on port 53088.
16/01/30 08:16:37 INFO netty.NettyBlockTransferService: Server created on 53088
16/01/30 08:16:37 INFO storage.BlockManagerMaster: Trying to register BlockManager
16/01/30 08:16:37 INFO storage.BlockManagerMasterEndpoint: Registering block manager 192.168.1.121:53088 with 517.4 MB RAM, BlockManagerId(driver, 192.168.1.121, 53088)
16/01/30 08:16:37 INFO storage.BlockManagerMaster: Registered BlockManager
16/01/30 08:16:38 INFO client.AppClient\$ClientEndpoint: Executor updated: app-20160130081636-0001/0 is now RUNNING
16/01/30 08:16:38 INFO client.AppClient\$ClientEndpoint: Executor updated: app-20160130081636-0001/1 is now RUNNING
16/01/30 08:16:38 INFO client.AppClient\$ClientEndpoint: Executor updated: app-20160130081636-0001/2 is now RUNNING
16/01/30 08:16:43 INFO cluster.SparkDeploySchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minRegisteredResourcesRatio: 0.0
16/01/30 08:17:00 INFO storage.MemoryStore: Block broadcast_0 stored as values in memory (estimated size 127.8 KB, free 127.8 KB)
16/01/30 08:17:01 INFO cluster.SparkDeploySchedulerBackend: Registered executor NettyRpcEndpointRef(null) (slq3:52579) with ID 0
16/01/30 08:17:02 INFO storage.BlockManagerMasterEndpoint: Registering block manager slq3:54589 with 517.4 MB RAM, BlockManagerId(0, slq3, 54589)
16/01/30 08:17:02 INFO storage.MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 14.3 KB, free 142.1 KB)
16/01/30 08:17:02 INFO cluster.SparkDeploySchedulerBackend: Registered executor NettyRpcEndpointRef(null) (slq2:52779) with ID 1
16/01/30 08:17:02 INFO storage.BlockManagerInfo: Added broadcast_0_piece0 in memory on 192.168.1.121:53088 (size: 14.3 KB, free: 517.4 MB)
16/01/30 08:17:03 INFO spark.SparkContext: Created broadcast 0 from textFile at WordCountCluster.scala:36
16/01/30 08:17:03 INFO storage.BlockManagerMasterEndpoint: Registering block manager slq2:38295 with 517.4 MB RAM, BlockManagerId(1, slq2, 38295)
16/01/30 08:17:23 INFO mapred.FileInputFormat: Total input paths to process : 1
16/01/30 08:17:26 INFO spark.SparkContext: Starting job: collect at WordCountCluster.scala:51
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Registering RDD 3 (map at WordCountCluster.scala:46)
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Got job 0 (collect at WordCountCluster.scala:51) with 2 output partitions
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Final stage: ResultStage 1 (collect at WordCountCluster.scala:51)
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Parents of final stage: List(ShuffleMapStage 0)
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Missing parents: List(ShuffleMapStage 0)
16/01/30 08:17:27 INFO scheduler.DAGScheduler: Submitting ShuffleMapStage 0 (MapPartitionsRDD[3] at map at WordCountCluster.scala:46), which has no missing parents
16/01/30 08:17:28 INFO storage.MemoryStore: Block broadcast_1 stored as values in memory (estimated size 4.1 KB, free 146.2 KB)
16/01/30 08:17:29 INFO storage.MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 2.3 KB, free 148.5 KB)
16/01/30 08:17:29 INFO storage.BlockManagerInfo: Added broadcast_1_piece0 in memory on 192.168.1.121:53088 (size: 2.3 KB, free: 517.4 MB)
16/01/30 08:17:29 INFO spark.SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1006
16/01/30 08:17:29 INFO scheduler.DAGScheduler: Submitting 2 missing tasks from ShuffleMapStage 0 (MapPartitionsRDD[3] at map at WordCountCluster.scala:46)
16/01/30 08:17:29 INFO scheduler.TaskSchedulerImpl: Adding task set 0.0 with 2 tasks
16/01/30 08:17:30 INFO scheduler.TaskSetManager: Starting task 0.0 in stage 0.0 (TID 0, slq3, partition 0,NODE_LOCAL, 2192 bytes)
16/01/30 08:17:30 INFO scheduler.TaskSetManager: Starting task 1.0 in stage 0.0 (TID 1, slq2, partition 1,NODE_LOCAL, 2192 bytes)
16/01/30 08:17:34 INFO storage.BlockManagerInfo: Added broadcast_1_piece0 in memory on slq2:38295 (size: 2.3 KB, free: 517.4 MB)
16/01/30 08:17:34 INFO storage.BlockManagerInfo: Added broadcast_1_piece0 in memory on slq3:54589 (size: 2.3 KB, free: 517.4 MB)
16/01/30 08:17:36 INFO storage.BlockManagerInfo: Added broadcast_0_piece0 in memory on slq2:38295 (size: 14.3 KB, free: 517.4 MB)
16/01/30 08:17:37 INFO storage.BlockManagerInfo: Added broadcast_0_piece0 in memory on slq3:54589 (size: 14.3 KB, free: 517.4 MB)
16/01/30 08:17:47 INFO scheduler.TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 18118 ms on slq3 (1/2)
16/01/30 08:17:47 INFO scheduler.TaskSetManager: Finished task 1.0 in stage 0.0 (TID 1) in 17770 ms on slq2 (2/2)
16/01/30 08:17:48 INFO scheduler.DAGScheduler: ShuffleMapStage 0 (map at WordCountCluster.scala:46) finished in 18.399 s
16/01/30 08:17:48 INFO scheduler.TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool
16/01/30 08:17:48 INFO scheduler.DAGScheduler: looking for newly runnable stages
16/01/30 08:17:48 INFO scheduler.DAGScheduler: running: Set()
16/01/30 08:17:48 INFO scheduler.DAGScheduler: waiting: Set(ResultStage 1)
16/01/30 08:17:48 INFO scheduler.DAGScheduler: failed: Set()
16/01/30 08:17:48 INFO scheduler.DAGScheduler: Submitting ResultStage 1 (ShuffledRDD[4] at reduceByKey at WordCountCluster.scala:50), which has no missing parents
16/01/30 08:17:48 INFO storage.MemoryStore: Block broadcast_2 stored as values in memory (estimated size 2.6 KB, free 151.1 KB)
16/01/30 08:17:48 INFO storage.MemoryStore: Block broadcast_2_piece0 stored as bytes in memory (estimated size 1593.0 B, free 152.6 KB)
16/01/30 08:17:48 INFO storage.BlockManagerInfo: Added broadcast_2_piece0 in memory on 192.168.1.121:53088 (size: 1593.0 B, free: 517.4 MB)
16/01/30 08:17:48 INFO spark.SparkContext: Created broadcast 2 from broadcast at DAGScheduler.scala:1006
16/01/30 08:17:48 INFO scheduler.DAGScheduler: Submitting 2 missing tasks from ResultStage 1 (ShuffledRDD[4] at reduceByKey at WordCountCluster.scala:50)
16/01/30 08:17:48 INFO scheduler.TaskSchedulerImpl: Adding task set 1.0 with 2 tasks
16/01/30 08:17:48 INFO scheduler.TaskSetManager: Starting task 0.0 in stage 1.0 (TID 2, slq3, partition 0,NODE_LOCAL, 1949 bytes)

16/01/30 08:17:48 INFO scheduler.TaskSetManager: Starting task 1.0 in stage 1.0 (TID 3, slq2, partition 1,NODE_LOCAL, 1949 bytes)

16/01/30 08:17:49 INFO storage.BlockManagerInfo: Added broadcast_2_piece0 in memory on slq2:38295 (size: 1593.0 B, free: 517.4 MB)

16/01/30 08:17:49 INFO storage.BlockManagerInfo: Added broadcast_2_piece0 in memory on slq3:54589 (size: 1593.0 B, free: 517.4 MB)

16/01/30 08:17:49 INFO spark.MapOutputTrackerMasterEndpoint: Asked to send map output locations for shuffle 0 to slq2:52779

16/01/30 08:17:49 INFO spark.MapOutputTrackerMaster: Size of output statuses for shuffle 0 is 151 bytes

16/01/30 08:17:49 INFO spark.MapOutputTrackerMasterEndpoint: Asked to send map output locations for shuffle 0 to slq3:52579

16/01/30 08:17:51 INFO scheduler.TaskSetManager: Finished task 1.0 in stage 1.0 (TID 3) in 2404 ms on slq2 (1/2)

16/01/30 08:17:51 INFO scheduler.TaskSetManager: Finished task 0.0 in stage 1.0 (TID 2) in 2531 ms on slq3 (2/2)

16/01/30 08:17:51 INFO scheduler.TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool

16/01/30 08:17:51 INFO scheduler.DAGScheduler: ResultStage 1 (collect at WordCountCluster.scala:51) finished in 2.538 s

16/01/30 08:17:51 INFO scheduler.DAGScheduler: Job 0 finished: collect at WordCountCluster.scala:51, took 24.676312 s

package : 1

this : 1

Version"](<http://spark.apache.org/docs/latest/building-spark.html#specifying-the-hadoop-version>) : 1

Because : 1

Python : 2

cluster. : 1

its : 1

[run : 1

general : 2

have : 1

pre-built : 1

locally. : 1

locally : 2

changed : 1

sc.parallelize(1 : 1

only : 1

several : 1

This : 2

basic : 1

Configuration : 1

learning. : 1

documentation : 3

YARN, : 1

graph : 1

Hive : 2

first : 1

[\"Specifying : 1

\"yarn-client\" : 1

page](<http://spark.apache.org/documentation.html>) : 1

[params]'. : 1

application : 1

[project : 2

prefer : 1

SparkPi : 2

<<http://spark.apache.org/>> : 1

engine : 1

version : 1

file : 1

documentation, : 1

MASTER : 1

example : 3

distribution. : 1

are : 1

params : 1

scala> : 1

systems. : 1

provides : 1

refer : 2

configure : 1

Interactive : 2

can : 6

build : 3

when : 1

easiest : 1

Apache : 1

Distributions"](<http://spark.apache.org/docs/latest/hadoop-third-party-distributions.html>) : 1

works : 1

how : 2

package. : 1

1000).count() : 1

Note : 1

Data. : 1

>>> : 1

Scala : 2

Alternatively, : 1

variable : 1

submit : 1

Testing : 1

Streaming : 1

thread. : 1

rich : 1

them, : 1

detailed : 2

stream : 1

GraphX : 1

tests](<https://cwiki.apache.org/confluence/display/SPARK/Contributing+to+Spark#ContributingtoSpark-AutomatedTesting>). : 1

distribution : 1

Please : 3

return : 2

is : 6

Thriftserver : 1

[\"Third : 1

same : 1

start : 1

built : 1

one : 2

with : 4

Party : 1

Spark](#building-spark). : 1

Spark"](<http://spark.apache.org/docs/latest/building-spark.html>). : 1

data : 2

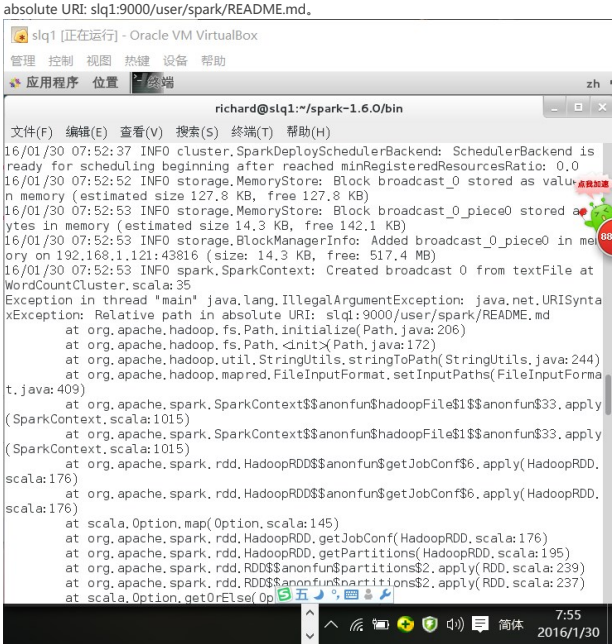
wiki](<https://cwiki.apache.org/confluence/display/SPARK>). : 1

using : 2

automated : 1
talk : 1
Shell : 2
class : 2
README : 1
computing : 1
Python, : 2
example : 1
: 8
from : 1
set : 2
building : 3
N : 1
Hadoop-supported : 1
other : 1
Example : 1
analysis. : 1
runs. : 1
Building : 1
higher-level : 1
need : 1
Big : 1
fast : 1
guide, : 1
Java, : 1
<class> : 1
uses : 1
SQL : 2
will : 1
guidance : 3
requires : 1
: 66
Documentation : 1
web : 1
cluster : 2
using : 1
MLlib : 1
shell : 2
Scala, : 1
supports : 2
built, : 1
./dev/run-tests : 1
sample : 1
For : 2
Spark : 14
particular : 3
Programs : 1
The : 1
processing. : 1
APIs : 1
computation : 1
Try : 1
[Configuration : 1
library : 1
A : 1
through : 1
: 1
./bin/pyspark : 1
following : 2
"yarn-cluster" : 1
More : 1
which : 2
See : 1
also : 5
storage : 1
for : 11
should : 2
To : 2
Once : 1
mesos:// : 1
setup : 1
Maven](<http://maven.apache.org/>). : 1
your : 1
latest : 1
processing, : 2
the : 21
not : 1
different : 1
guide](<http://spark.apache.org/docs/latest/configuration.html>) : 1
distributions. : 1
given. : 1
About : 1
if : 4
instructions. : 1
be : 2
Tests : 1
do : 2
no : 1
all : 1
./bin/run-example : 2
programs, : 1
including : 3
./bin/run-example : 1
Spark. : 1
Versions : 1
HDFS : 1
spark:// : 1
It : 2
an : 3
programming : 1
machine : 1
environment : 1
run : 1
clean : 1
1000 : 2
And : 1
run : 7
URL, : 1
./bin/spark-shell : 1
threads. : 1
"local" : 1
MASTER=spark://host:7077 : 1
on : 6

```
You : 3
against : 1
help : 1
print : 1
tests : 1
[Apache : 1
examples : 2
at : 2
in : 5
-DskipTests : 1
usage : 1
downloaded : 1
versions : 1
online : 1
graphs : 1
optimized : 1
abbreviated : 1
comes : 1
directory. : 1
overview : 1
[building : 1
`examples` : 2
Many : 1
Running : 1
way : 1
use : 3
Online : 1
site, : 1
running : 1
find : 1
sc.parallelize(range(1000)).count() : 1
contains : 1
you : 4
project : 1
name : 1
protocols : 1
that : 3
a : 9
or : 3
high-level : 1
Pi : 1
Hadoop, : 2
to : 14
available : 1
core : 1
(You : 1
instance: : 1
see : 1
of : 5
tools : 1
"local[N]" : 1
programs : 2
structured : 1
package.): 1
["Building : 1
must : 1
and : 10
command, : 2
system : 1
mvn : 1
Hadoop : 4
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/metrics/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/stage/kill,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/api,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/static,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped
o.s.j.s.ServletContextHandler(/executors/threadDump/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped
o.s.j.s.ServletContextHandler(/executors/threadDump,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/executors/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/executors,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/environment/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/environment,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/storage/rdd/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/storage/rdd,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/storage/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/storage,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/pool/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/pool,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/pool/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/stage/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/stage,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/stages,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/jobs/job/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/jobs/job,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/jobs/json,null)
16/01/30 08:17:52 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHandler(/jobs,null)
16/01/30 08:17:52 INFO ui.SparkUI: Stopped Spark web UI at http://192.168.1.121:4040
16/01/30 08:17:52 INFO cluster.SparkDeploySchedulerBackend: Shutting down all executors
16/01/30 08:17:52 INFO cluster.SparkDeploySchedulerBackend: Asking each executor to shut down
16/01/30 08:17:52 INFO spark.MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
16/01/30 08:17:53 INFO storage.MemoryStore: MemoryStore cleared
16/01/30 08:17:53 INFO storage.BlockManager: BlockManager stopped
16/01/30 08:17:53 INFO storage.BlockManagerMaster: BlockManagerMaster stopped
16/01/30 08:17:53 INFO scheduler.OutputCommitCoordinator$OutputCommitCoordinatorEndpoint:
OutputCommitCoordinator stopped!
16/01/30 08:17:53 INFO spark.SparkContext: Successfully stopped SparkContext
16/01/30 08:17:53 INFO remote.RemoteActorRefProvider$RemotingTerminator: Shutting down remote daemon.
16/01/30 08:17:53 INFO remote.RemoteActorRefProvider$RemotingTerminator: Remote daemon shut down;
proceeding with flushing remote transports.
16/01/30 08:17:54 INFO util.ShutdownHookManager: Shutdown hook called
16/01/30 08:17:54 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-42b0d08b-938a-45fe-8dae-
81db69953855/httpd-ccd8a47c-1b7f-4ac6-b7e9-510b4f88c4b2
16/01/30 08:17:54 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-42b0d08b-938a-45fe-8dae-
81db69953855
16/01/30 08:17:55 INFO remote.RemoteActorRefProvider$RemotingTerminator: Remoting shut down.
[richard@slq1 bin]$
```

小插曲：
第一次运行时报错：
Exception in thread "main" java.lang.IllegalArgumentException: java.net.URISyntaxException: Relative path in



原因：
val lines = sc.textFile("hdfs://192.168.1.121:9000/user/spark/README.md") 原来写成相对路
径：/user/spark/README.md，导致程序找不到文件。

以上内容是王家林老师DT大数据梦工厂《IMF传奇行动》第8课的学习笔记。
王家林: Spark, Flink, Docker, Android技术中国区布道师。Spark亚太研究院院长和首席专家, DT大数据梦工厂创始人, Android软硬整合源码
级专家, 英语发音魔术师, 健身狂热爱好者。
微信公众号: DT_Spark
联系邮箱18610086859@126.com
电话: 18610086859
QQ: 1740415547
微信号: 18610086859
新浪微博: ilovepains
王家林老师的第一个中国梦: 免费为全社会培养100万名优秀的大数据从业人员！

版权声明：本文为博主原创文章，未经博主允许不得转载。

目前您尚未登录, 请 [登录](#) 或 [注册](#) 后进行评论

相关文章推荐

spark-submit

[jifeng@feng03 spark-1.4.0-bin-hadoop2.6]\$./bin/spark-submit --class "SimpleApp" --master spark://f...

wind520 2015年08月20日 23:28 2179

Flume环境部署和配置详解及案例大全

flume是一个分布式、可靠、和高可用的海量日志采集、聚合和传输的系统。支持在日志系统中定制各类数据发送方，用于收集数据;同时，Flu
me提供对数据进行简单处理，并写到各种数据接受方(比如文本、HDF...

moonpure 2015年07月22日 13:36 460

Spark 入门实战之最好的实例

转载：https://www.ibm.com/developerworks/cn/opensource/os-cn-spark-practice1/ 搭建开发环境 安装 Scala L...

sysmedia 2017年04月06日 17:18 372

第8课：彻底实战详解使用IDE开发Spark程序

本地安装好java和scala并创建工程WordCount

Diamonds_y5 2016年08月25日 16:35 258

第8课：彻底实战详解使用IDE开发Spark程序

第8课：彻底实战详解使用IDE开发Spark程序 1.下载安装windows下的scala-2.10.4. 2.打开eclipse，新建scala project：WordCount 3.修改
依赖的...

slq1023 2016年01月24日 00:13 2523

3000门徒内部训练绝密视频（泄密版）第8课：彻底实战详解使用IDE开发Spark程序

彻底实战详解使用IDE开发Spark程序使用IDE开发Spark分析 使用IDE开发Spark实战 使用IDE开发Spark的Local和Cluster开发两种选择：IDE
A、Eclipse下载S...

tom_8899_li 2016年08月11日 21:55 178

Scala-IDE Eclipse (Windows)中开发Spark应用程序，在Ubuntu Spark集群上运行

Scala-IDE Eclipse(Windows)中开发Spark应用程序，在Ubuntu Spark集群上运行在进行实际的Spark应用程序开发时，常常会利用Window
环境进行程序开发，开发测试好...

lovehuangjiaju 2015年07月11日 18:43 5167

大数据Spark “蘑菇云”行动前传第3课：在IDE下开发第一个Scala程序透彻解析及Scala控制结构详解实战

大数据Spark “蘑菇云”行动前传第3课：在IDE下开发第一个Scala程序透彻解析及Scala控制结构详解实战 1 通过IDE 开发第一个scala程序初

底详解 2 scala基础语法和控制结构实...

 duan_zhihua 2016年07月20日 20:53 568

Local模式下开发第一个Spark程序并运行于集群环境

第一阶段（1-3月）：会从浅入深，基于大量案例实战，深度剖析和讲解Spark，并且会包含完全从企业真实复杂业务需求中抽取出的案例实战。课程会涵盖Scala编程详解、Spark核心编程、Spark SQ...

 kxr0502 2016年01月10日 08:09 1268

Spark on yarn IntelliJ ide 安装，编译，打包，集群运行 详解

—intellij 安装centos65系统 步骤一 步骤二 步骤三本地运行 打包集群 yarn 运行说明：已经安装好hadoop2.2.0 完全分布，scala，spark已安装好，环境配置完毕；主...

 ZHAOLEI5911 2017年03月28日 14:26 388

第8课：Spark Streaming源码解读之RDD生成全生命周期彻底研究和思考

本期内容 1. DStream与RDD关系彻底研究 2. Streaming中RDD的生成彻底研究

 chinsun_1 2016-06-01 18:34 272

第7课在自己搭建的spark集群下运行程序实例

第一阶段：Spark streaming、spark sql、kafka、spark内核原理（必须有一个大型项目经验）； 第二阶段：spark运行的各种环境，各种故障的解决，性能优化（精通spark...

 lhui798 2016-01-10 19:53 288

第8课：Spark Streaming源码解读之RDD生成全生命周期彻底研究和思考

本期内容 Dstream与rdd关系彻底研究 Dstream中rdd生成彻底研究从整个sparkstreaming角度来考虑，rdd分为三个方面的内容： a.怎么生成的，依靠什么生成的； b...

 MichaelLi916 2016-05-29 00:59 1487

【Spark亚太研究院系列丛书】Spark实战高手之路-第一章 构建Spark集群-配置Hadoop-伪分布模式并运...

执行文件拷贝操作

 wwttz1974 2014-08-27 11:17 657

第8课：零基础实战Scala最常用数据结构Map和Tuple及Spark源码鉴赏

零基础实战Scala最常用数据结构Map和Tuple及Spark源码鉴赏

 w517424787 2016-07-30 21:34 892

第10课：底实战详解使用Java开发Spark程序学习笔记（二）

Maven下的Spark配置：http://maven.outofmemory.cn/org.apache.spark，这个网站提供了Spark core、Spark Streaming使用Mave...

 slq1023 2016-02-28 10:15 1564

第3课：在IDE下开发第一个Scala程序纯傻瓜式彻底透彻解析

Scala零基础实战详解！

 w517424787 2016-07-30 09:26 1673

第10课：底实战详解使用Java开发Spark程序学习笔记

第10课：底实战详解使用Java开发Spark程序学习笔记 本期内容： 1. 为什么要使用Java？ 2. 使用Java开发Spark实战 3. 使用Java开发Spark的Local和Cluster...

 slq1023 2016-02-27 00:00 1056

第95讲:使用Scala开发集群运行的Spark来实现在线热搜词条获取

package com.dt.streaming import org.apache.spark.SparkConf import org.apache.spark.streaming.(Secon...

 qq_21234493 2016-05-15 16:37 564

spark mllib 应用程序开发及提交到spark集群运行--入门

一、程序开发 Dependencies MLib uses the linear algebra package Breeze, which depends on netlib-java f...

 hualizhuanshen2014 2016-03-24 13:16 544